

```
In [1]: # This Python 3 environment comes with many helpful analytics libraries installed
# It is defined by the kaggle/python Docker image: https://github.com/kaggle/docker-python
# For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/) that gets preserved as output when you create a version using
# You can also write temporary files to /kaggle/temp/, but they won't be saved outside of the current session
```

/kaggle/input/my-data/Customer Loyalty History.csv  
/kaggle/input/my-data/Airline Loyalty Data Dictionary.csv  
/kaggle/input/my-data/Customer Flight Activity.csv

## Customer loyalty program data from Northern Lights Air (NLA) in Canada Analysis and Customer Lifetime Value(CLV) Prediction

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
import warnings
warnings.filterwarnings('ignore')
pd.set_option('display.max_columns', 50)
```

```
In [3]: df_activity=pd.read_csv('/kaggle/input/my-data/Customer Flight Activity.csv')
df_loyalty=pd.read_csv('/kaggle/input/my-data/Customer Loyalty History.csv')
```

```
In [4]: df_activity.head(3)
```

```
Out[4]:
```

	Loyalty Number	Year	Month	Flights Booked	Flights with Companions	Total Flights	Distance	Points Accumulated	Points Redeemed	Dollar Cost Points Redeemed
0	100018	2017	1	3	0	3	1521	152.0	0	0
1	100102	2017	1	10	4	14	2030	203.0	0	0
2	100140	2017	1	6	0	6	1200	120.0	0	0

```
In [5]: df_loyalty.head(3)
```

```
Out[5]:
```

	Loyalty Number	Country	Province	City	Postal Code	Gender	Education	Salary	Marital Status	Loyalty Card	CLV	Enrollment Type	Enrollment Year	Enrollment Month	Cancellation Year	Canc
0	480934	Canada	Ontario	Toronto	M2Z 4K1	Female	Bachelor	83236.0	Married	Star	3839.14	Standard	2016	2	NaN	
1	549612	Canada	Alberta	Edmonton	T3G 6Y6	Male	College	NaN	Divorced	Star	3839.61	Standard	2016	3	NaN	
2	429460	Canada	British Columbia	Vancouver	V6E 3D9	Male	College	NaN	Single	Star	3839.75	Standard	2014	7	2018.0	

## Merging all data

```
In [6]: df=df_activity.merge(df_loyalty,on='Loyalty Number')
```

In [7]: df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 405624 entries, 0 to 405623
Data columns (total 25 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Loyalty Number                        405624 non-null  int64
1   Year                                405624 non-null  int64
2   Month                              405624 non-null  int64
3   Flights Booked                       405624 non-null  int64
4   Flights with Companions              405624 non-null  int64
5   Total Flights                       405624 non-null  int64
6   Distance                             405624 non-null  int64
7   Points Accumulated                   405624 non-null  float64
8   Points Redeemed                     405624 non-null  int64
9   Dollar Cost Points Redeemed          405624 non-null  int64
10  Country                             405624 non-null  object
11  Province                            405624 non-null  object
12  City                                405624 non-null  object
13  Postal Code                          405624 non-null  object
14  Gender                              405624 non-null  object
15  Education                           405624 non-null  object
16  Salary                              302952 non-null  float64
17  Marital Status                      405624 non-null  object
18  Loyalty Card                        405624 non-null  object
19  CLV                                 405624 non-null  float64
20  Enrollment Type                     405624 non-null  object
21  Enrollment Year                     405624 non-null  int64
22  Enrollment Month                    405624 non-null  int64
23  Cancellation Year                    50064 non-null   float64
24  Cancellation Month                  50064 non-null   float64
dtypes: float64(5), int64(11), object(9)
memory usage: 77.4+ MB

```

In [8]: df.isna().sum()

```

Out[8]: Loyalty Number      0
Year      0
Month      0
Flights Booked      0
Flights with Companions      0
Total Flights      0
Distance      0
Points Accumulated      0
Points Redeemed      0
Dollar Cost Points Redeemed      0
Country      0
Province      0
City      0
Postal Code      0
Gender      0
Education      0
Salary      102672
Marital Status      0
Loyalty Card      0
CLV      0
Enrollment Type      0
Enrollment Year      0
Enrollment Month      0
Cancellation Year      355560
Cancellation Month      355560
dtype: int64

```

**Dropping missing values in Cancellation Year& Cancellation& Month&Postal Code and Country because the missing values are more than 85% from all data**

In [9]: df=df.drop(['Cancellation Year','Cancellation Month','Postal Code','Country'],axis=1)

**Completing missing values in Salary column by mean from simple imputer technique**

```
In [10]: from sklearn.impute import SimpleImputer

# Create an instance of SimpleImputer with the desired strategy
imputer = SimpleImputer(strategy='mean')

# Fit the imputer on the salary column and transform the missing values
df['Salary'] = imputer.fit_transform(df[['Salary']])
```

```
In [11]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 405624 entries, 0 to 405623
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Loyalty Number                        405624 non-null  int64
1   Year                                  405624 non-null  int64
2   Month                                405624 non-null  int64
3   Flights Booked                        405624 non-null  int64
4   Flights with Companions               405624 non-null  int64
5   Total Flights                         405624 non-null  int64
6   Distance                              405624 non-null  int64
7   Points Accumulated                    405624 non-null  float64
8   Points Redeemed                       405624 non-null  int64
9   Dollar Cost Points Redeemed           405624 non-null  int64
10  Province                              405624 non-null  object
11  City                                  405624 non-null  object
12  Gender                                405624 non-null  object
13  Education                             405624 non-null  object
14  Salary                                405624 non-null  float64
15  Marital Status                        405624 non-null  object
16  Loyalty Card                          405624 non-null  object
17  CLV                                   405624 non-null  float64
18  Enrollment Type                       405624 non-null  object
19  Enrollment Year                       405624 non-null  int64
20  Enrollment Month                      405624 non-null  int64
dtypes: float64(3), int64(11), object(7)
memory usage: 65.0+ MB
```

```
In [12]: df.head()
```

```
Out[12]:
```

	Loyalty Number	Year	Month	Flights Booked	Flights with Companions	Total Flights	Distance	Points Accumulated	Points Redeemed	Dollar Cost Points Redeemed	Province	City	Gender	Education	Salary	Marital Status
0	100018	2017	1	3	0	3	1521	152.0	0	0	Alberta	Edmonton	Female	Bachelor	92552.0	Married
1	100018	2017	2	2	2	4	1320	132.0	0	0	Alberta	Edmonton	Female	Bachelor	92552.0	Married
2	100018	2018	10	6	4	10	3110	311.0	385	31	Alberta	Edmonton	Female	Bachelor	92552.0	Married
3	100018	2017	4	4	0	4	924	92.0	0	0	Alberta	Edmonton	Female	Bachelor	92552.0	Married
4	100018	2017	5	0	0	0	0	0.0	0	0	Alberta	Edmonton	Female	Bachelor	92552.0	Married

## Extracting top 10 customers having the highest average CLV

```
In [14]: df.groupby('Loyalty Number')['CLV'].mean().sort_values(ascending=False).head(10)
```

```
Out[14]: Loyalty Number
615459      83325.38
652627      83325.38
776187      74228.52
844145      74228.52
767366      73225.96
592003      73225.96
838263      67907.27
680886      67907.27
179870      66025.75
495253      66025.75
Name: CLV, dtype: float64
```

## Extracting top 10 customers having the highest Total Flights number


```
In [15]: df.groupby('Loyalty Number')['Total Flights'].sum().sort_values(ascending=False).head(10)
```

```
Out[15]: Loyalty Number
336882    448
775768    400
464187    399
255836    395
512296    392
736504    383
279419    379
215508    375
876062    372
615561    372
Name: Total Flights, dtype: int64
```

## Exploratory Data Analysis(EDA)


```
In [16]: df['City'].value_counts().plot(kind='bar')
```

```
Out[16]: <Axes: xlabel='City'>
```




```
In [17]: df['Gender'].value_counts().plot(kind='pie', autopct='%1.1f%%')
```

```
Out[17]: <Axes: ylabel='count'>
```



```
In [18]: df['Marital Status'].value_counts().plot(kind='pie', autopct='%1.1f%%')
```

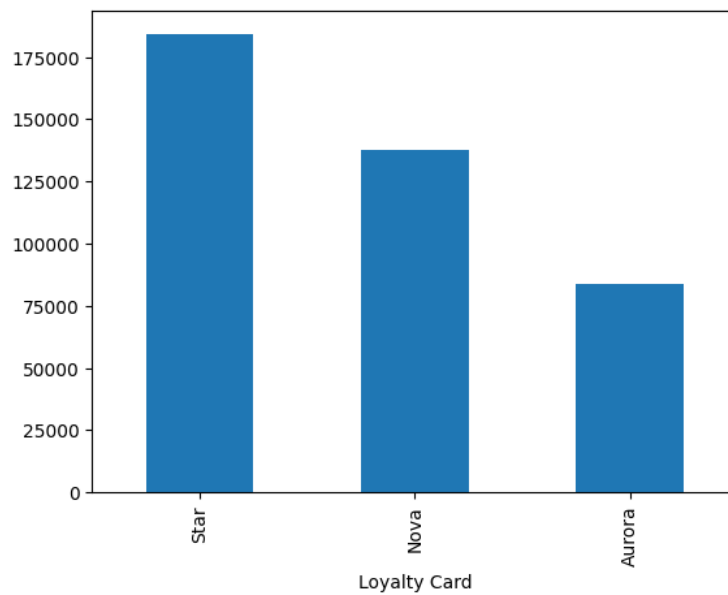
```
Out[18]: <Axes: ylabel='count'>
```



### Most customers are from married people

```
In [19]: df['Loyalty Card'].value_counts().plot(kind='bar')
```


```
Out[19]: <Axes: xlabel='Loyalty Card'>
```



### The highest selling loyalty card is star and the lowest is aurora

```
In [20]: df['Education'].value_counts().plot(kind='bar')
```


```
Out[20]: <Axes: xlabel='Education'>
```



## The most customers are having Bachelor degree

```
In [21]: df['Enrollment Type'].value_counts().plot(kind='pie', autopct='%1.1f%%')
```

```
Out[21]: <Axes: ylabel='count'>
```



```
In [22]: df=df.drop('Loyalty Number',axis=1)
```

```
In [23]: df['Enrollment Type'].value_counts()
```

```
Out[23]: Enrollment Type
Standard      382200
2018 Promotion  23424
Name: count, dtype: int64
```

## Converting Gender and Enrollment Type columns to numerical columns

```
In [24]: df['Gender']=df['Gender'].replace({'Female':0,'Male':1})
df['Enrollment Type']=df['Enrollment Type'].replace({'Standard':0,'2018 Promotion':1})
```

## Converting categorical features to numerical by LabelEncoder to can use them in prediction.

```
In [25]: from sklearn.preprocessing import LabelEncoder
labelencoder=LabelEncoder()
df['Province']=labelencoder.fit_transform(df['Province'])
df['City']=labelencoder.fit_transform(df['City'])
df['Education']=labelencoder.fit_transform(df['Education'])
df['Marital Status']=labelencoder.fit_transform(df['Marital Status'])
df['Loyalty Card']=labelencoder.fit_transform(df['Loyalty Card'])
```

## Extracting the correlations between all features and CLV column to determine the features that have the highest influence on CLV feature.

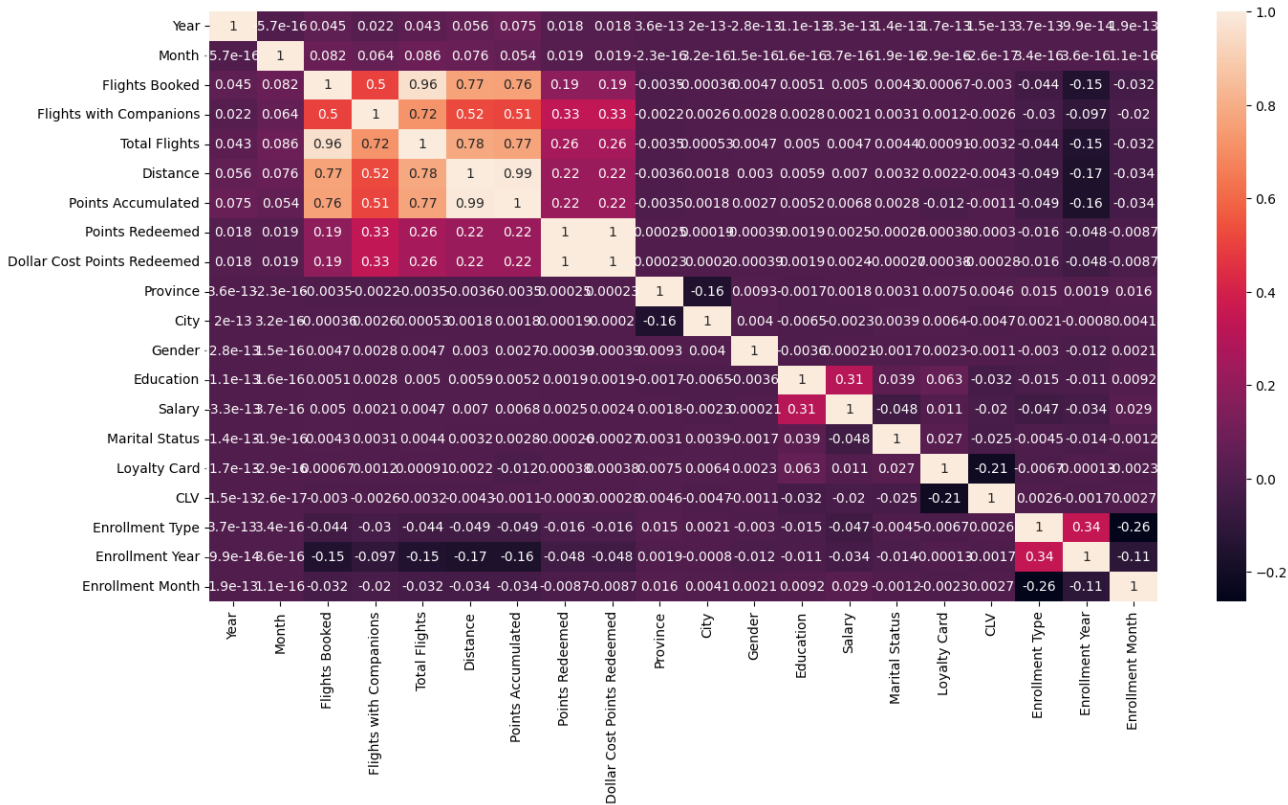
```
In [26]: df.corr()
```

```
Out[26]:
```

	Year	Month	Flights Booked	Flights with Companions	Total Flights	Distance	Points Accumulated	Points Redeemed	Dollar Cost Points Redeemed	Province	City	CLV
Year	1.000000e+00	-5.737962e-16	0.044510	0.021615	0.042550	0.056140	0.075208	0.017633	0.017624	3.580244e-13	2.043782e-13	-2.83
Month	-5.737962e-16	1.000000e+00	0.082133	0.064492	0.086353	0.076345	0.054178	0.019408	0.019315	-2.307385e-16	3.152263e-16	1.4965
Flights Booked	4.451043e-02	8.213330e-02	1.000000	0.502500	0.961344	0.767457	0.760279	0.188232	0.188242	-3.461029e-03	-3.605528e-04	4.7096
Flights with Companions	2.161493e-02	6.449206e-02	0.502500	1.000000	0.721136	0.517979	0.511090	0.334325	0.334385	-2.193267e-03	2.562142e-03	2.8248
Total Flights	4.255050e-02	8.635348e-02	0.961344	0.721136	1.000000	0.779935	0.771989	0.257307	0.257334	-3.471862e-03	5.270724e-04	4.6735
Distance	5.614031e-02	7.634481e-02	0.767457	0.517979	0.779935	1.000000	0.994564	0.224261	0.224283	-3.625086e-03	1.791207e-03	3.0077
Points Accumulated	7.520848e-02	5.417842e-02	0.760279	0.511090	0.771989	0.994564	1.000000	0.223230	0.223264	-3.455821e-03	1.759084e-03	2.7329
Points Redeemed	1.763278e-02	1.940776e-02	0.188232	0.334325	0.257307	0.224261	0.223230	1.000000	0.999972	2.510058e-04	1.888936e-04	-3.86
Dollar Cost Points Redeemed	1.762408e-02	1.931472e-02	0.188242	0.334385	0.257334	0.224283	0.223264	0.999972	1.000000	2.325534e-04	1.974144e-04	-3.92
Province	3.580244e-13	-2.307385e-16	-0.003461	-0.002193	-0.003472	-0.003625	-0.003456	0.000251	0.000233	1.000000e+00	-1.558457e-01	9.2722
City	2.043782e-13	3.152263e-16	-0.000361	0.002562	0.000527	0.001791	0.001759	0.000189	0.000197	-1.558457e-01	1.000000e+00	4.0353
Gender	-2.834385e-13	1.496595e-16	0.004710	0.002825	0.004674	0.003008	0.002733	-0.000387	-0.000392	9.272205e-03	4.035342e-03	1.0000
Education	-1.086689e-13	1.626831e-16	0.005135	0.002838	0.005018	0.005853	0.005157	0.001889	0.001902	-1.741454e-03	-6.454917e-03	-3.62
Salary	-3.284742e-13	3.736994e-16	0.005010	0.002150	0.004699	0.006951	0.006802	0.002450	0.002445	1.843328e-03	-2.266529e-03	2.1481
Marital Status	-1.434437e-13	-1.916412e-16	0.004303	0.003115	0.004440	0.003229	0.002805	-0.000265	-0.000274	3.133414e-03	3.881080e-03	-1.66
Loyalty Card	-1.707973e-13	-2.915448e-16	0.000669	0.001160	0.000905	0.002201	-0.012249	0.000382	0.000379	7.527903e-03	6.350424e-03	2.2982
CLV	1.488269e-13	-2.556763e-17	-0.002964	-0.002583	-0.003197	-0.004252	-0.001075	-0.000304	-0.000277	4.567617e-03	-4.713820e-03	-1.10
Enrollment Type	3.697842e-13	3.445648e-16	-0.043759	-0.029551	-0.044476	-0.049261	-0.049275	-0.016238	-0.016254	1.521648e-02	2.096196e-03	-3.04
Enrollment Year	-9.927540e-14	3.564005e-16	-0.153455	-0.097012	-0.153861	-0.165263	-0.164645	-0.047867	-0.047915	1.909355e-03	-7.986071e-04	-1.23
Enrollment Month	1.890369e-13	1.100554e-16	-0.031528	-0.019644	-0.031520	-0.034416	-0.033893	-0.008689	-0.008692	1.589009e-02	4.096361e-03	2.0552

```
In [27]: plt.figure(figsize=(16,8))
sns.heatmap(df.corr(),annot=True)
```

Out[27]: <Axes: >



```
In [28]: df.head()
```

Out[28]:

	Year	Month	Flights Booked	Flights with Companions	Total Flights	Distance	Points Accumulated	Points Redeemed	Dollar Cost Points Redeemed	Province	City	Gender	Education	Salary	Marital Status	Loyalty Card	CLV
0	2017	1	3	0	3	1521	152.0	0	0	0	4	0	0	92552.0	1	0	7919
1	2017	2	2	2	4	1320	132.0	0	0	0	4	0	0	92552.0	1	0	7919
2	2018	10	6	4	10	3110	311.0	385	31	0	4	0	0	92552.0	1	0	7919
3	2017	4	4	0	4	924	92.0	0	0	0	4	0	0	92552.0	1	0	7919
4	2017	5	0	0	0	0	0.0	0	0	0	4	0	0	92552.0	1	0	7919

```
In [29]: X=df.drop('CLV',axis=1)
y=df.CLV
```

## Splitting the data to train and test

```
In [30]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=.25,stratify=y,random_state=42)
```

## Scaling all data by MinMaxScaler to increase the prediction accuracy



```
In [31]: from sklearn.preprocessing import MinMaxScaler
scale=MinMaxScaler()
X_train_scaled=scale.fit_transform(X_train)
X_train_scaled
```

```
Out[31]: array([[0.          , 0.          , 0.33333333, ..., 0.          , 0.66666667,
        1.          ],
       [0.          , 0.72727273, 0.38095238, ..., 0.          , 0.83333333,
        0.09090909],
       [1.          , 1.          , 0.28571429, ..., 0.          , 0.5          ,
        0.36363636],
       ...,
       [0.          , 0.90909091, 0.33333333, ..., 0.          , 0.83333333,
        0.27272727],
       [1.          , 1.          , 0.23809524, ..., 0.          , 0.5          ,
        1.          ],
       [1.          , 0.81818182, 0.          , ..., 0.          , 0.16666667,
        0.45454545]])
```

```
In [32]: X_test_scaled=scale.transform(X_test)
X_test_scaled
```

```
Out[32]: array([[1.          , 0.          , 0.          , ..., 1.          , 1.          ,
        0.27272727],
       [1.          , 0.09090909, 0.38095238, ..., 0.          , 0.          ,
        0.45454545],
       [1.          , 0.54545455, 0.23809524, ..., 0.          , 1.          ,
        0.          ],
       ...,
       [0.          , 0.72727273, 0.61904762, ..., 0.          , 0.66666667,
        1.          ],
       [0.          , 0.63636364, 0.          , ..., 0.          , 1.          ,
        0.54545455],
       [0.          , 1.          , 0.04761905, ..., 0.          , 0.16666667,
        0.27272727]])
```

```
In [33]: from sklearn.metrics import mean_absolute_error, mean_squared_error
```

## Using random forest regressor for CLV prediction

```
In [34]: from sklearn.ensemble import RandomForestRegressor
RFR=RandomForestRegressor(n_estimators=300)
RFR.fit(X_train_scaled,y_train)
y_pred=RFR.predict(X_test_scaled)
print(mean_absolute_error(y_pred,y_test))
```

```
270.3215078472331
```

## Using XGB Regressor for CLV prediction

```
In [35]: from xgboost import XGBRegressor
XGB=XGBRegressor(n_estimators=1000,learning_rate=1,max_depth=13)
XGB.fit(X_train_scaled,y_train)
y_pred2=XGB.predict(X_test_scaled)
print(mean_absolute_error(y_pred2,y_test))
```

```
849.2057255379643
```

```
In [36]: y_pred[:10]
```

```
Out[36]: array([ 3115.7          , 3530.23383333, 15943.436          , 4621.43349167,
        11638.9          , 33806.78743333, 5674.93586667, 6288.7598          ,
        2916.64466667, 8622.84          ])
```

```
In [37]: y_test.head(10)
```

```
Out[37]: 245412    3115.70
356677    3472.37
343052    16702.70
137897    4766.37
2895      11638.90
186527    34090.04
395808    5660.13
29135     6301.10
77628     2722.21
287876    8622.84
Name: CLV, dtype: float64
```

## Comparing between actual and prediction CLV to determine the accuracy and

```
In [38]: pred_df=pd.DataFrame(y_pred,index=X_test.index,columns=['prediction'])
# Merge the predicted values with the actual test data
merged_df = pd.concat([y_test,pred_df], axis=1)
merged_df=merged_df.rename(columns={'CLV':'Actual'})
merged_df
```

Out[38]:

	Actual	prediction
245412	3115.70	3115.700000
356677	3472.37	3530.233833
343052	16702.70	15943.436000
137897	4766.37	4621.433492
2895	11638.90	11638.900000
...	...	...
74002	3097.92	3130.740467
269131	3919.37	3919.370000
281504	3237.50	3290.053233
28471	8236.92	8236.920000
366563	17759.39	8745.155633

101406 rows × 2 columns

**As we see the random forest regressor algorithm is the best one for CLV prediction**